

Basic ML Guide

This notebook covers developing and tuning simple Classification and Regression machine learning models in python.

- Dataset used for analysis is [Sarah Gets a Diamond](#) taken from University of Virginia, Darden Business Publishing.
- The regression model predicts the price(Y variable) of a diamond based on its multiple physical attributes/features(X variables). For the classification model, I have changed the definition of the Y variable. For all diamonds with price greater than \$10,000 High/H label was added and for the remaining Low/L label was added.
- Basic data cleaning, data transformations, and pre-processing techniques have been applied on the on model development data. Model development and Hyper-parameter tuning has been carried out for both the Regression and Classification models.

In [1]:

```
import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
import shap

import xgboost as xgb
import lightgbm as lgb

from scipy.stats import gaussian_kde
from sklearn.svm import SVR, SVC
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder, MinMaxScaler
from sklearn.model_selection import KFold, cross_validate, train_test_split, cross_val_score, RandomizedSearchCV
from sklearn.utils import check_array

from sklearn.metrics import make_scorer, r2_score, mean_absolute_error, roc_curve, auc, classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, precision_score, recall_score, accuracy_score

import warnings

warnings.filterwarnings('ignore')
pd.set_option('display.float_format', lambda x : '%.5f' % x)
```

pandas.Int64Index is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

1. Defining utility functions to be used later

In [2]: *# function to prepare Confusion Matrix, RoC-AUC curve, and other relvant statistics for a Classification problem*

```
def clf_report(clf, x_true, y_true, split):
    y_pred = clf.predict(x_true)
    probs = clf.predict_proba(x_true)
    print('Classification report for {} data'.format(split))
    cm = confusion_matrix(y_true, y_pred, labels=clf.classes_)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
    disp.plot()
    plt.show()
    print('Overall Accuracy : {}'.format(round(accuracy_score(y_true, y_pred) * 100, 2)))
    print('Precision Score : {}'.format(round(precision_score(y_true, y_pred, average='binary') * 100, 2)))
    print('Recall Score : {}'.format(round(recall_score(y_true, y_pred, average='binary') * 100, 2)))
    preds = probs[:,1]
    fpr, tpr, threshold = roc_curve(y_true, preds)
    roc_auc = auc(fpr, tpr)
    print('AUC : {}'.format(round(roc_auc * 100, 2)))
    plt.figure()
    plt.plot(fpr, tpr, label='AUC = %0.2f' % roc_auc)
    plt.plot([0.0, 1.0], [0, 1], 'r--')
    plt.xlim([-0.1, 1.1])
    plt.ylim([-0.1, 1.1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('RoC-AUC on {} Data'.format(split))
    plt.legend(loc="lower right")
    plt.show()
    print('\n')
```

In [3]: *# functions to calculate MAPE and Negative-MAPE for Regression problems*

```
def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

def neg_mean_absolute_percentage_error(y_true, y_pred):
    return (-1)*(mean_absolute_percentage_error(y_true, y_pred))
```

In [4]: *# function to check a Regression model performanace on Train data*

```

# Performance measured on R2, MAE, MAPE metrics on 10-fold Cross-Validation

def train_metrics(rgs, X_cv, Y_cv):
    print('Cross Validated Metric Results for Train Data:')
    kf = KFold(n_splits=10, random_state=25, shuffle=True)
    metric_df = pd.DataFrame()

    for train_index, test_index in kf.split(X_cv):
        x_train, x_test = X_cv.iloc[train_index], X_cv.iloc[test_index]
        y_train, y_test = Y_cv.iloc[train_index], Y_cv.iloc[test_index]
        y_pred = rgs.fit(x_train, y_train).predict(x_test)
        r2 = r2_score(y_test, y_pred)
        mae = mean_absolute_error(y_test, y_pred)
        mape = mean_absolute_percentage_error(y_test, y_pred)
        temp_df = pd.DataFrame([[r2, mae, mape]])
        metric_df = pd.concat([metric_df, temp_df], ignore_index = True)

    metric_df.loc['Mean'] = round((metric_df.mean()),2)
    metric_df.loc['Std Dev'] = round((metric_df.std()),2)
    metric_df = metric_df.set_axis(['R-sq', 'MAE', 'MAPE'], axis=1, inplace=False)

    with pd.option_context('float_format', '{:.2f}'.format, 'display.expand_frame_repr', False):
        print(metric_df, '\n')

```

In [5]:

```

# function to check a Regression model performanace on Test data
# Performance measured on R2, MAE, MAPE metrics on 10-fold Cross-Validation

def test_metrics(Y_test, Y_pred):
    print('Metric Results for Test Data:')
    r2 = r2_score(Y_test, Y_pred)
    mae = mean_absolute_error(Y_test, Y_pred)
    mape = mean_absolute_percentage_error(Y_test, Y_pred)
    metric_df = pd.DataFrame([[r2, mae, mape]], columns=['R-sq', 'MAE', 'MAPE'])
    with pd.option_context('float_format', '{:.2f}'.format, 'display.expand_frame_repr', False):
        print(metric_df, '\n')

```

2. Data import

In [6]:

```

df_full = pd.read_csv('./sgd.csv')
display(df_full.shape)

```

```
display(df_full.head())
display(df_full.tail())
```

(9142, 9)

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	5169.00000
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	3470.00000
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	3183.00000
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	4370.00000
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	3171.00000

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
9137	9138	0.96000	Ideal	F	SI1	EX	EX	GIA	NaN
9138	9139	1.02000	Very Good	E	VVS1	EX	G	GIA	NaN
9139	9140	1.51000	Good	I	VS1	G	G	GIA	NaN
9140	9141	1.24000	Ideal	H	VS2	VG	VG	GIA	NaN
9141	9142	0.79000	Ideal	I	VS1	EX	EX	GIA	NaN

In [7]:

```
# selecting the rows with non-null values for the target/Y variable
```

```
df = df_full.copy()
df = df[:6000]
display(df.shape)
display(df.head())
display(df.tail())
```

(6000, 9)

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	5169.00000
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	3470.00000
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	3183.00000

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	4370.00000
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	3171.00000

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
5995	5996	1.03000	Ideal	D	SI1	EX	EX	GIA	6250.00000
5996	5997	1.00000	Very Good	D	SI1	VG	VG	GIA	5328.00000
5997	5998	1.02000	Ideal	D	SI1	EX	EX	GIA	6157.00000
5998	5999	1.27000	Signature-Ideal	G	VS1	EX	EX	GIA	11206.00000
5999	6000	2.19000	Ideal	E	VS1	EX	EX	GIA	30507.00000

```
In [8]: # descriptive statistics for categorical variables
df.describe(include='object')
```

```
Out[8]:
```

	Cut	Color	Clarity	Polish	Symmetry	Report
count	6000	6000	6000	6000	6000	6000
unique	5	6	7	4	4	2
top	Ideal	G	SI1	EX	VG	GIA
freq	2482	1501	2059	2425	2417	5266

```
In [9]: # descriptive statistics for numeric variables across 3 terciles
df.describe(include='number', percentiles=[0.33,0.66])
```

```
Out[9]:
```


	ID	Carat Weight	Price
count	6000.00000	6000.00000	6000.00000
mean	3000.50000	1.33452	11791.57933
std	1732.19514	0.47570	10184.35005

	ID	Carat Weight	Price
min	1.00000	0.75000	2184.00000
33%	1980.67000	1.02000	5789.00000
50%	3000.50000	1.13000	7857.00000
66%	3960.34000	1.50000	11083.36000
max	6000.00000	2.91000	101561.00000

```
In [10]: # checking distribution for Price variable.
# This will help in creating a new target variable definition for the classification problem

df['Price'].plot.hist()
```

```
Out[10]: <AxesSubplot:ylabel='Frequency'>
```



```
In [11]: count = df[df['Price'] < 10000].count()
print(count)
```

```
ID          3657
Carat Weight 3657
Cut          3657
Color        3657
Clarity      3657
```

```
Polish          3657
Symmetry        3657
Report          3657
Price           3657
dtype: int64
```

```
In [12]: # creating a new variable/feature - Rate. This will serve as the target variable for the classification problem

df['Rate'] = np.where(((df['Price'] <= 10000)), 'L', 'H')
display(df.head())
```

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price	Rate
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	5169.00000	L
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	3470.00000	L
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	3183.00000	L
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	4370.00000	L
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	3171.00000	L

```
In [13]: df['Rate'].value_counts(), df['Rate'].value_counts().sum()
```

```
Out[13]: (L      3658
         H      2342
         Name: Rate, dtype: int64,
         6000)
```

```
In [14]: # dataset for the classification problem

df_clf = df.copy()
df_clf = df_clf.drop('Price', axis=1)
df_clf.head()
```

```
Out[14]:
```

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	L
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	L
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	L

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	L
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	L

```
In [15]: # dataset for the regression problem

df_reg = df.copy()
df_reg = df_reg.drop('Rate', axis=1)
df_reg.head()
```

```
Out[15]:
```

	ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	5169.00000
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	3470.00000
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	3183.00000
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	4370.00000
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	3171.00000

3. Classification

3.1 Data Cleaning and Exploratory Analysis

```
In [16]: # descriptive statistics for categorical variables

df_clf.describe(include='object')
```

```
Out[16]:
```

	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
count	6000	6000	6000	6000	6000	6000	6000
unique	5	6	7	4	4	2	2
top	Ideal	G	SI1	EX	VG	GIA	L
freq	2482	1501	2059	2425	2417	5266	3658


```
In [17]: # descriptive statistics for numeric variables

df_clf.describe(include='number')
```

```
Out[17]:
```

	ID	Carat Weight
count	6000.00000	6000.00000
mean	3000.50000	1.33452
std	1732.19514	0.47570
min	1.00000	0.75000
25%	1500.75000	1.00000
50%	3000.50000	1.13000
75%	4500.25000	1.59000
max	6000.00000	2.91000

```
In [18]: # removing white spaces from the column names

dictionary = {' ': '_', '-': ''}
df_clf.replace(dictionary, regex=True, inplace=True)
df_clf.columns = df_clf.columns.str.replace(' ', '_')
display(df_clf.head())
display(df_clf.tail())
```

	ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	L
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	L
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	L
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	L
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	L


	ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
--	----	--------------	-----	-------	---------	--------	----------	--------	------

	ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Rate
5995	5996	1.03000	Ideal	D	SI1	EX	EX	GIA	L
5996	5997	1.00000	Very_Good	D	SI1	VG	VG	GIA	L
5997	5998	1.02000	Ideal	D	SI1	EX	EX	GIA	L
5998	5999	1.27000	SignatureIdeal	G	VS1	EX	EX	GIA	H
5999	6000	2.19000	Ideal	E	VS1	EX	EX	GIA	H

In [19]:

```
# Target Class/Variable distribution : Count and Percentage wise

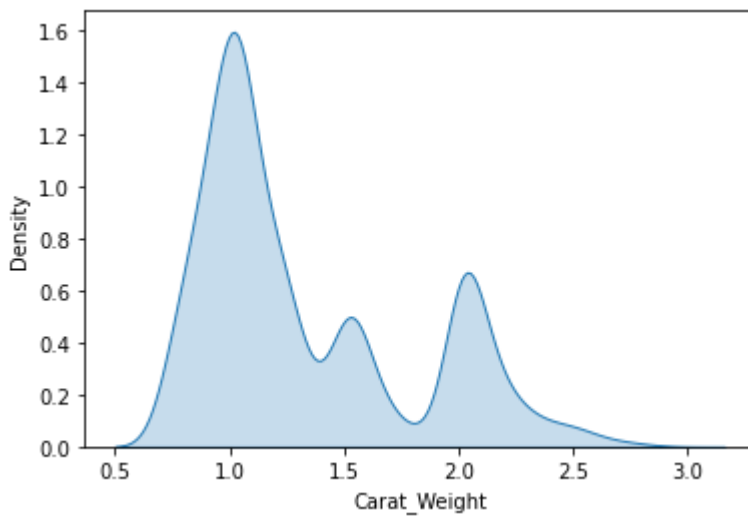
grid = sns.catplot(x='Rate', kind='count', data=df_clf)
ax = grid.axes[0, 0]
ax.bar_label(ax.containers[0])
for p in ax.patches:
    percentage = '{}%'.format(round(100 * (p.get_height()) / (len(df_clf))),2)
    x = p.get_x() + (p.get_width())/2
    y = p.get_height() - 0.05 * (p.get_height())
    ax.annotate(percentage, (x, y), ha='center')
plt.show()
```



In [20]:


```
# density plot for Carat_Weight to see if some mathematical transformations are needed for this variable
```

```
sns.kdeplot(df_clf['Carat_Weight'], fill='bool')  
plt.show()
```




```
In [21]: # density plot for Carat_Weight across different Target classes. High Rate diamonds tend to have higher carat weight
```


```
sns.kdeplot(data=df_clf, x='Carat_Weight', hue='Rate', fill='bool')  
plt.show()
```



```
In [22]: # histogram plots with a different visualization library to check Carat_Weight distribution
```

```
fig = px.histogram(df_clf, x='Carat_Weight', histnorm='probability density')  
fig.show()
```






In [23]:

histogram plots with a different visualization library to check Carat_Weight distribution across different classes

```
fig = px.histogram(df_clf, x='Carat_Weight', color='Rate')
fig.show()
```



3.2 Pre-processing and Feature Engineering

In [24]: *# creating new features with some mathematical-transformation*

```
df_clf['log_CW'] = np.log(df_clf['Carat_Weight'])
df_clf['norm_CW'] = MinMaxScaler().fit_transform(df_clf[['Carat_Weight']])
df_clf['lgm_CW'] = np.log((df_clf['norm_CW']+0.00001))
```

In [25]: df_clf.dtypes

Out[25]:

ID	int64
Carat_Weight	float64
Cut	object
Color	object
Clarity	object
Polish	object
Symmetry	object
Report	object
Rate	object
log_CW	float64
norm_CW	float64
lgm_CW	float64
dtype:	object

In [26]: *# collecting all the numerical and categorical variables and saving them in two different lists*

```
categorical_columns_seen = [c for i,c in enumerate(df_clf.columns) if df_clf.dtypes[i] in [object]]
categorical_columns_seen.remove('Rate')
numerical_columns_seen = [c for i,c in enumerate(df_clf.columns) if df_clf.dtypes[i] not in [object]]
categorical_columns_seen, numerical_columns_seen
```


```
Out[26]: (['Cut', 'Color', 'Clarity', 'Polish', 'Symmetry', 'Report'],  
          ['ID', 'Carat_Weight', 'log_CW', 'norm_CW', 'lgnm_CW'])
```

```
In [27]: # creating one hot encoded labels for categorical data  
  
encoder = OneHotEncoder(handle_unknown='error', drop='first')  
encoded_df = (pd.DataFrame(encoder.fit_transform(df_clf[categorical_columns_seen]).toarray()).astype(int)  
encoded_df.columns = encoder.get_feature_names_out(categorical_columns_seen)  
encoded_df.head()
```

```
Out[27]:
```

	Cut_Good	Cut_Ideal	Cut_SignatureIdeal	Cut_Very_Good	Color_E	Color_F	Color_G	Color_H	Color_I	Clarity_IF	...	Clarity_VS2	Clarity_VVS
0	0	1	0	0	0	0	0	1	0	0	...	0	
1	0	1	0	0	0	0	0	1	0	0	...	0	
2	0	1	0	0	0	0	0	1	0	0	...	0	
3	0	1	0	0	1	0	0	0	0	0	...	0	
4	0	1	0	0	0	0	1	0	0	0	...	0	

5 rows × 22 columns



```
In [28]: tmp_df = pd.DataFrame()  
tmp_df[numerical_columns_seen] = df_clf[numerical_columns_seen]  
tmp_df.head()
```

```
Out[28]:
```

	ID	Carat_Weight	log_CW	norm_CW	lgnm_CW
0	1	1.10000	0.09531	0.16204	-1.81987
1	2	0.83000	-0.18633	0.03704	-3.29557
2	3	0.85000	-0.16252	0.04630	-3.07248
3	4	0.91000	-0.09431	0.07407	-2.60255
4	5	0.83000	-0.18633	0.03704	-3.29557

```
In [29]: # converting Target Variable/Class into 1-0
```

```
cols = list(encoded_df.columns.values)
tmp_df[cols] = encoded_df[cols]
tmp_df['Rate'] = np.where(((df_clf['Rate'] == 'H')), 1, 0)
tmp_df.head()
```

Out[29]:

	ID	Carat_Weight	log_CW	norm_CW	lgnm_CW	Cut_Good	Cut_Ideal	Cut_SignatureIdeal	Cut_Very_Good	Color_E	...	Clarity_VVS1	Clarity
0	1	1.10000	0.09531	0.16204	-1.81987	0	1	0	0	0	...	0	
1	2	0.83000	-0.18633	0.03704	-3.29557	0	1	0	0	0	...	0	
2	3	0.85000	-0.16252	0.04630	-3.07248	0	1	0	0	0	...	0	
3	4	0.91000	-0.09431	0.07407	-2.60255	0	1	0	0	1	...	0	
4	5	0.83000	-0.18633	0.03704	-3.29557	0	1	0	0	0	...	0	

5 rows × 28 columns



3.3 Model Development (feature set -1)

In [30]:

```
# creating Train and test data

Y_clf = tmp_df['Rate']
X_clf = tmp_df.drop(['Rate', 'ID', 'log_CW', 'norm_CW', 'lgnm_CW'], axis = 1)

X_train_clf, X_test_clf, Y_train_clf, Y_test_clf = train_test_split(X_clf, Y_clf, train_size = 0.8334, random_state = 25)
X_train_clf.shape, X_test_clf.shape
```

Out[30]: ((5000, 23), (1000, 23))

In [31]:

```
shap.initjs()
```



In [32]:

```
# model 1 : Light Gradient Boosting model

lgb_model = lgb.LGBMClassifier(boosting_type='goss', num_leaves=16, max_depth=-1, learning_rate=0.1, n_estimators=64, n_j
```




```
lgb_model.fit(X_train_clf, Y_train_clf, objective='binary', reg_lambda=0.5, importance_type='gain', silent=True, random_state=25)
```


```
Out[32]: LGBMClassifier(boosting_type='goss', importance_type='gain', n_estimators=64,  
num_leaves=16, objective='binary', random_state=25,  
reg_lambda=0.5)
```

```
In [33]: # checking model performance on train and test data  
  
clf_report(lgb_model, X_train_clf, Y_train_clf, 'Train')  
clf_report(lgb_model, X_test_clf, Y_test_clf, 'Test')
```


Classification report for Train data




Overall Accuracy : 98.34
Precision Score : 97.58
Recall Score : 98.14
AUC : 99.91



Classification report for Test data



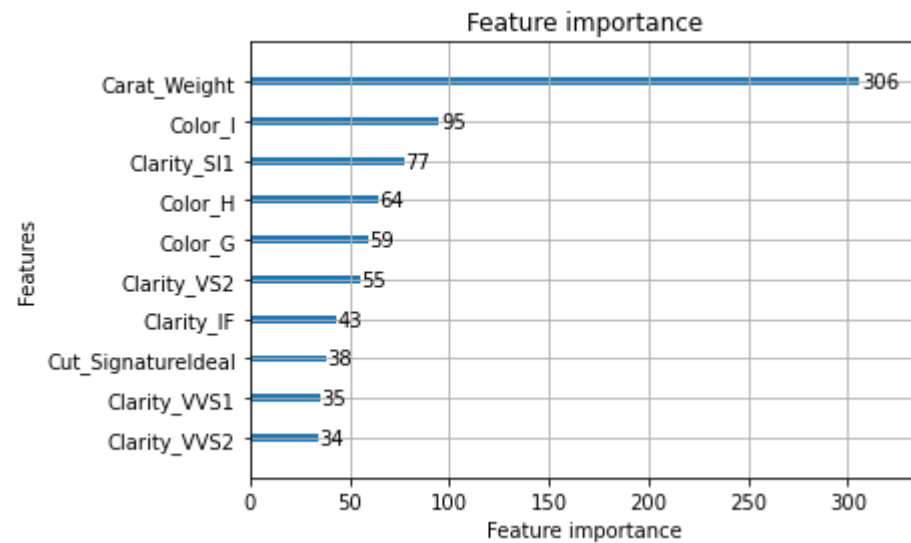
Overall Accuracy : 96.5
Precision Score : 96.07
Recall Score : 95.37
AUC : 99.52



```
In [34]: # creating feature importance plot

lgb.plot_importance(lgb_model, max_num_features=10)
```


```
Out[34]: <AxesSubplot:title={'center':'Feature importance'}, xlabel='Feature importance', ylabel='Features'>
```



```
In [35]:
```

```
# creating SHAP-explainer plot for all features
```


```
explainer = shap.TreeExplainer(lgb_model)
shap_values = explainer.shap_values(X_train_clf)
shap.summary_plot(shap_values[1], X_train_clf, max_display=10)
plt.show()
```



```
In [36]: # creating SHAP-explainer plot for all features but for a single data row
```

```
shap.force_plot(explainer.expected_value[1], shap_values[1][0,:], X_train_clf.iloc[0, :])
```

```
Out[36]:
```



```
In [37]: # model 2 : Xtreme Gradient Boosting model


xgb_model = xgb.XGBClassifier(n_estimators=64, max_depth=6, learning_rate=0.1, verbosity=0, use_label_encoder=False,
                             booster='gbtree', n_jobs=-1, reg_lambda=0.3, random_state=25)
xgb_model.fit(X_train_clf, Y_train_clf)
```

```
Out[37]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                    gamma=0, gpu_id=-1, importance_type=None,
                    interaction_constraints='', learning_rate=0.1, max_delta_step=0,
                    max_depth=6, min_child_weight=1, missing=nan,
                    monotone_constraints='()', n_estimators=64, n_jobs=-1,
                    num_parallel_tree=1, predictor='auto', random_state=25,
                    reg_alpha=0, reg_lambda=0.3, scale_pos_weight=1, subsample=1,
                    tree_method='exact', use_label_encoder=False,
                    validate_parameters=1, verbosity=0)
```


```
In [38]: # checking model performance on Train and test data

clf_report(xgb_model, X_train_clf, Y_train_clf, 'Train')
clf_report(xgb_model, X_test_clf, Y_test_clf, 'Test')
```


Classification report for Train data




Overall Accuracy : 98.76
Precision Score : 98.3
Recall Score : 98.5
AUC : 99.94



Classification report for Test data




Overall Accuracy : 96.3
Precision Score : 96.05
Recall Score : 94.88
AUC : 99.47



```
In [39]: # creating feature importance plot

xgb.plot_importance(xgb_model, max_num_features=10)
```

```
Out[39]: <AxesSubplot:title={'center':'Feature importance'}, xlabel='F score', ylabel='Features'>
```



```
In [40]:
```


```
# creating features tree plot to explain xgb-model decisions
```

```
xgb_params = xgb_model.get_booster()
for importance_type in ('weight', 'gain', 'cover', 'total_gain', 'total_cover'):
    print('%s: ' % importance_type, xgb_params.get_score(importance_type=importance_type))

xgb.to_graphviz(xgb_model, numtrees=0, size='20,20')
```

```
weight: {'Carat_Weight': 540.0, 'Cut_Good': 16.0, 'Cut_Ideal': 52.0, 'Cut_SignatureIdeal': 82.0, 'Cut_Very_Good': 50.0,
'Color_E': 31.0, 'Color_F': 59.0, 'Color_G': 109.0, 'Color_H': 117.0, 'Color_I': 118.0, 'Clarity_IF': 79.0, 'Clarity_SI
1': 133.0, 'Clarity_VS1': 44.0, 'Clarity_VS2': 72.0, 'Clarity_VVS1': 61.0, 'Clarity_VVS2': 57.0, 'Polish_G': 47.0, 'Polis
h_ID': 26.0, 'Polish_VG': 34.0, 'Symmetry_G': 27.0, 'Symmetry_ID': 8.0, 'Symmetry_VG': 48.0, 'Report_GIA': 10.0}
gain: {'Carat_Weight': 46.63850402832031, 'Cut_Good': 1.3568055629730225, 'Cut_Ideal': 2.0487358570098877, 'Cut_Signatur
eIdeal': 3.1732118129730225, 'Cut_Very_Good': 1.293332576751709, 'Color_E': 4.7530059814453125, 'Color_F': 3.070007085800
171, 'Color_G': 5.6871466636657715, 'Color_H': 8.126121520996094, 'Color_I': 13.49828815460205, 'Clarity_IF': 6.334937572
479248, 'Clarity_SI1': 6.782624244689941, 'Clarity_VS1': 6.442461013793945, 'Clarity_VS2': 9.109304428100586, 'Clarity_VV
S1': 4.946499824523926, 'Clarity_VVS2': 6.01350736618042, 'Polish_G': 2.2005317211151123, 'Polish_ID': 1.005779027938842
8, 'Polish_VG': 1.5062283277511597, 'Symmetry_G': 1.3638572692871094, 'Symmetry_ID': 3.322267532348633, 'Symmetry_VG': 1.
0099114179611206, 'Report_GIA': 1.8627281188964844}
cover: {'Carat_Weight': 113.99594116210938, 'Cut_Good': 18.982723236083984, 'Cut_Ideal': 23.309589385986328, 'Cut_Signat
ureIdeal': 102.09420013427734, 'Cut_Very_Good': 10.433252334594727, 'Color_E': 24.423532485961914, 'Color_F': 21.39283943
1762695, 'Color_G': 34.65987014770508, 'Color_H': 41.618221282958984, 'Color_I': 58.61147689819336, 'Clarity_IF': 137.178
1768798828, 'Clarity_SI1': 50.38148880004883, 'Clarity_VS1': 36.72146224975586, 'Clarity_VS2': 52.17066955566406, 'Clarit
y_VVS1': 148.55996704101562, 'Clarity_VVS2': 164.23777770996094, 'Polish_G': 28.90943717956543, 'Polish_ID': 8.9628982543
94531, 'Polish_VG': 17.54803466796875, 'Symmetry_G': 23.14409065246582, 'Symmetry_ID': 19.828092575073242, 'Symmetry_VG':
7.35402250289917, 'Report_GIA': 6.514835357666016}
total_gain: {'Carat_Weight': 25184.79296875, 'Cut_Good': 21.70888900756836, 'Cut_Ideal': 106.53426361083984, 'Cut_Signat
ureIdeal': 260.203369140625, 'Cut_Very_Good': 64.6666259765625, 'Color_E': 147.3431854248047, 'Color_F': 181.130416870117
2, 'Color_G': 619.8989868164062, 'Color_H': 950.7562255859375, 'Color_I': 1592.7979736328125, 'Clarity_IF': 500.460083007
8125, 'Clarity_SI1': 902.0890502929688, 'Clarity_VS1': 283.4682922363281, 'Clarity_VS2': 655.8699340820312, 'Clarity_VVS
1': 301.7364807128906, 'Clarity_VVS2': 342.7699279785156, 'Polish_G': 103.42498779296875, 'Polish_ID': 26.15025520324707,
'Polish_VG': 51.211761474609375, 'Symmetry_G': 36.82414627075195, 'Symmetry_ID': 26.578140258789062, 'Symmetry_VG': 48.47
574996948242, 'Report_GIA': 18.627281188964844}
total_cover: {'Carat_Weight': 61557.80859375, 'Cut_Good': 303.72357177734375, 'Cut_Ideal': 1212.0986328125, 'Cut_Signatu
reIdeal': 8371.724609375, 'Cut_Very_Good': 521.66259765625, 'Color_E': 757.1295166015625, 'Color_F': 1262.177490234375,
'Color_G': 3777.92578125, 'Color_H': 4869.33203125, 'Color_I': 6916.154296875, 'Clarity_IF': 10837.076171875, 'Clarity_SI
1': 6700.73779296875, 'Clarity_VS1': 1615.744384765625, 'Clarity_VS2': 3756.2880859375, 'Clarity_VVS1': 9062.158203125,
'Clarity_VVS2': 9361.5537109375, 'Polish_G': 1358.7435302734375, 'Polish_ID': 233.0353546142578, 'Polish_VG': 596.6331787
109375, 'Symmetry_G': 624.8904418945312, 'Symmetry_ID': 158.62474060058594, 'Symmetry_VG': 352.9930725097656, 'Report_GI
A': 65.14835357666016}
```


Out[40]:



In [41]:

```
# creating SHAP-explainer plot for all features
```

```
explainer = shap.TreeExplainer(xgb_model)
shap_values = explainer.shap_values(X_train_clf)
shap.summary_plot(shap_values, X_train_clf, max_display=10)
```




In [42]:

```
# creating SHAP-explainer plot for all features but for a single data row
```

```
shap.force_plot(explainer.expected_value, shap_values[0,:], X_train_clf.iloc[0, :])
```

Out[42]:



3.4 Model Development (feature set-2)

```
In [43]: Y_clf = tmp_df['Rate']
X_clf = tmp_df.drop(['Rate', 'ID', 'lgnm_CW', 'log_CW', 'Carat_Weight'], axis = 1)

X_train_clf, X_test_clf, Y_train_clf, Y_test_clf = train_test_split(X_clf, Y_clf, train_size = 0.8334, random_state = 25)
X_train_clf.shape, X_test_clf.shape
```


```
Out[43]: ((5000, 23), (1000, 23))
```

```
In [44]: lgb_model = lgb.LGBMClassifier(boosting_type='goss', num_leaves=16, max_depth=-1, learning_rate=0.1, n_estimators=128, n
        objective='binary', reg_lambda=0.5, importance_type='gain', random_state=25, silent=True)
lgb_model.fit(X_train_clf, Y_train_clf)
```


```
Out[44]: LGBMClassifier(boosting_type='goss', importance_type='gain', n_estimators=128,
        num_leaves=16, objective='binary', random_state=25,
        reg_lambda=0.5)
```

```
In [45]: clf_report(lgb_model, X_train_clf, Y_train_clf, 'Train')
        clf_report(lgb_model, X_test_clf, Y_test_clf, 'Test')
```


Classification report for Train data



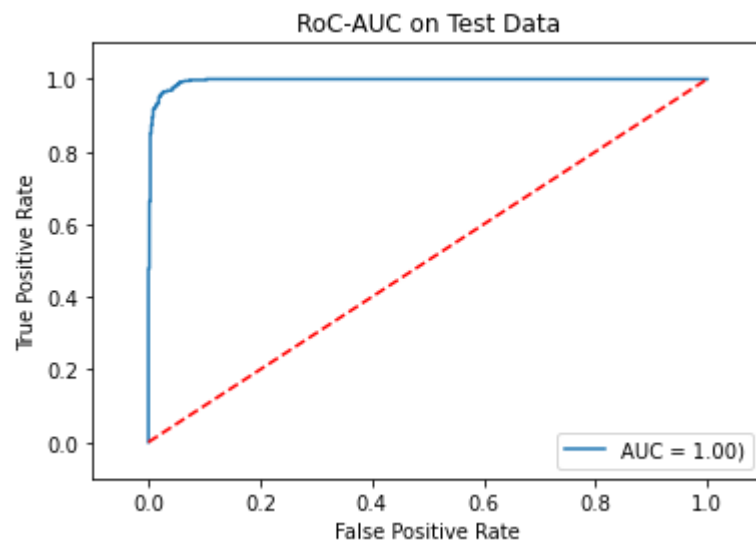
Overall Accuracy : 99.0
Precision Score : 98.61
Recall Score : 98.81
AUC : 99.97



Classification report for Test data



Overall Accuracy : 96.9
 Precision Score : 95.88
 Recall Score : 96.59
 AUC : 99.58



3.5 Hyperparameter Tuning

```
In [46]: params = {'n_estimators' : [64, 128, 256], 'max_depth' : [3,5,7], 'learning_rate' : [0.5, 0.1]}
randomized_object = RandomizedSearchCV(estimator=lgb.LGBMClassifier(), scoring='roc_auc', random_state=25,
```

```
cv=10, n_jobs=-1, param_distributions=params, verbose=0)
randomized_object.fit(X_train_clf, Y_train_clf)
```

```
Out[46]: RandomizedSearchCV(cv=10, estimator=LGBMClassifier(), n_jobs=-1,
                        param_distributions={'learning_rate': [0.5, 0.1],
                        'max_depth': [3, 5, 7],
                        'n_estimators': [64, 128, 256]},
                        random_state=25, scoring='roc_auc')
```

```
In [47]: # tuned hyperparameters

print('Best Parameters : {}'.format(randomized_object.best_params_))
print('Best AUC_score : {}'.format(round(randomized_object.best_score_, 4)))
print('Best model : {}'.format(randomized_object.best_estimator_))
CV_Res = pd.concat([pd.DataFrame(randomized_object.cv_results_['params']),
                    pd.DataFrame(randomized_object.cv_results_['mean_test_score'], columns=['AUC_score'])], axis=1)
CV_Res = CV_Res.sort_values(by='AUC_score', ascending=False)
print(CV_Res)
```

```
Best Parameters : {'n_estimators': 256, 'max_depth': 3, 'learning_rate': 0.1}
```

```
Best AUC_score : 0.9978
```


```
Best model : LGBMClassifier(max_depth=3, n_estimators=256)
```

	n_estimators	max_depth	learning_rate	AUC_score
0	256	3	0.10000	0.99782
4	64	3	0.50000	0.99776
6	128	3	0.50000	0.99744
2	128	3	0.10000	0.99737
3	128	7	0.10000	0.99736
1	256	5	0.10000	0.99726
5	64	5	0.50000	0.99705
9	256	7	0.10000	0.99701
7	256	5	0.50000	0.99636
8	64	3	0.10000	0.99551

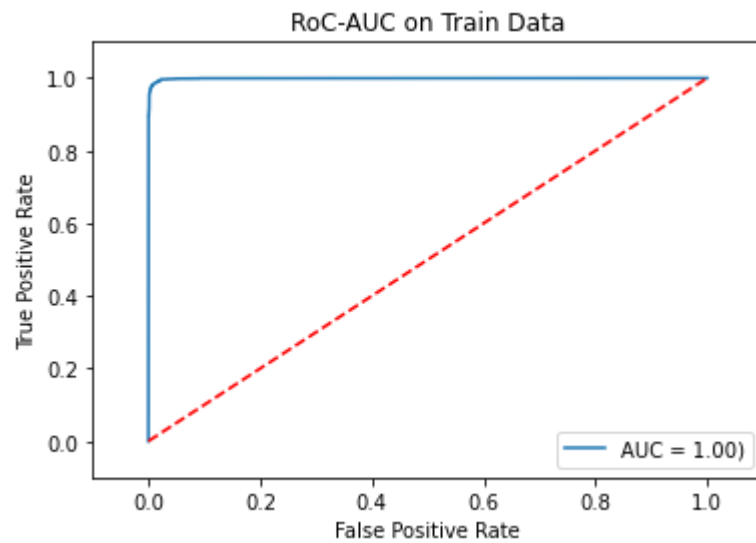
```
In [48]: # Tuned model performance on Train and Test data

clf_report(randomized_object.best_estimator_, X_train_clf, Y_train_clf, 'Train')
clf_report(randomized_object.best_estimator_, X_test_clf, Y_test_clf, 'Test')
```


Classification report for Train data




Overall Accuracy : 98.78
Precision Score : 98.25
Recall Score : 98.6
AUC : 99.94




Classification report for Test data




Overall Accuracy : 96.9
Precision Score : 96.56
Recall Score : 95.85
AUC : 99.6




```
In [49]: lgb.plot_importance(randomized_object.best_estimator_, max_num_features=10)
plt.show()
```



```
In [50]: explainer = shap.TreeExplainer(randomized_object.best_estimator_)
shap_values = explainer.shap_values(X_train_clf)
shap.summary_plot(shap_values[1], X_train_clf, max_display=10)
plt.show()
```

```
In [51]: shap.force_plot(explainer.expected_value[1], shap_values[1][0,:], X_train_clf.iloc[0, :])
```



4. Regression

```
In [52]: # descriptive statistics for categorical variables

df_reg.describe(include='object')
```

```
Out[52]:
```

Cut	Color	Clarity	Polish	Symmetry	Report
-----	-------	---------	--------	----------	--------

	Cut	Color	Clarity	Polish	Symmetry	Report
count	6000	6000	6000	6000	6000	6000
unique	5	6	7	4	4	2
top	Ideal	G	SI1	EX	VG	GIA
freq	2482	1501	2059	2425	2417	5266

```
In [53]: # descriptive statistics for numeric variables

df_reg.describe(include='number')
```

```
Out[53]:
```

	ID	Carat Weight	Price
count	6000.00000	6000.00000	6000.00000
mean	3000.50000	1.33452	11791.57933
std	1732.19514	0.47570	10184.35005
min	1.00000	0.75000	2184.00000
25%	1500.75000	1.00000	5150.50000
50%	3000.50000	1.13000	7857.00000
75%	4500.25000	1.59000	15036.50000
max	6000.00000	2.91000	101561.00000

4.1 Data Cleaning and Exploratory Analysis

```
In [54]: # removing white spaces from the column names

dictionary = {' ': '_', '-': ''}
df_reg.replace(dictionary, regex=True, inplace=True)
df_reg.columns = df_reg.columns.str.replace(' ', '_')
display(df_reg.head())
display(df_reg.tail())
```

ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
----	--------------	-----	-------	---------	--------	----------	--------	-------

	ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
0	1	1.10000	Ideal	H	SI1	VG	EX	GIA	5169.00000
1	2	0.83000	Ideal	H	VS1	ID	ID	AGSL	3470.00000
2	3	0.85000	Ideal	H	SI1	EX	EX	GIA	3183.00000
3	4	0.91000	Ideal	E	SI1	VG	VG	GIA	4370.00000
4	5	0.83000	Ideal	G	SI1	EX	EX	GIA	3171.00000

	ID	Carat_Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
5995	5996	1.03000	Ideal	D	SI1	EX	EX	GIA	6250.00000
5996	5997	1.00000	Very_Good	D	SI1	VG	VG	GIA	5328.00000
5997	5998	1.02000	Ideal	D	SI1	EX	EX	GIA	6157.00000
5998	5999	1.27000	SignatureIdeal	G	VS1	EX	EX	GIA	11206.00000
5999	6000	2.19000	Ideal	E	VS1	EX	EX	GIA	30507.00000

In [55]:

```
# creating new features with some mathematical-transformation


df_reg['log_CW'] = np.log(df_reg['Carat_Weight'])
df_reg['norm_CW'] = MinMaxScaler().fit_transform(df_reg[['Carat_Weight']])
df_reg['lgm_CW'] = np.log(df_reg['norm_CW']+0.00001)
```


In [56]:

```
# scatter plot for Carat_Weight with respect to different features to observe existing correlations

fig = px.scatter(df_reg, x='Carat_Weight', y='Price', color='Color', template='ggplot2')
fig.show()
```

100k—

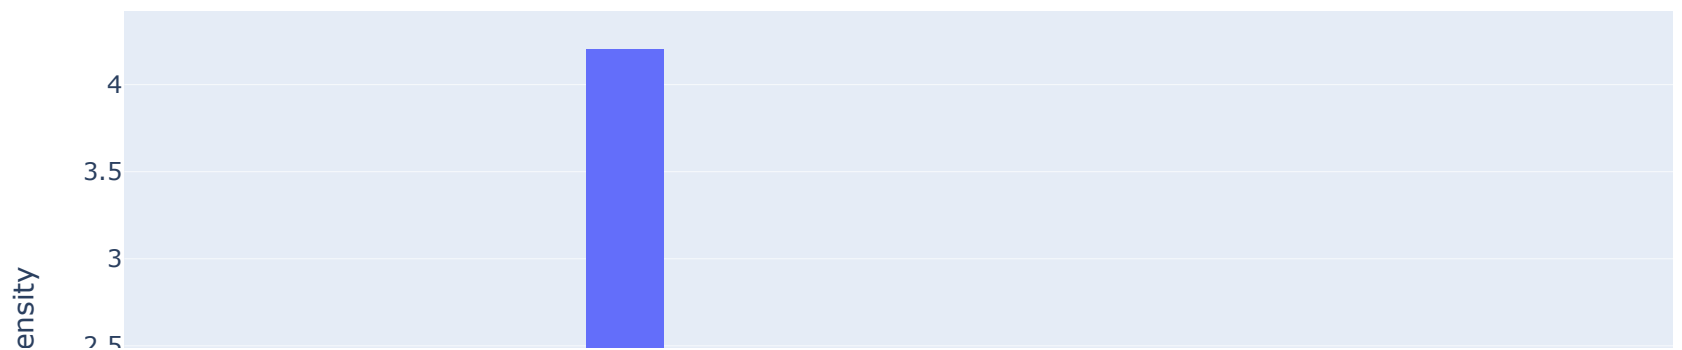





In [57]:

```
# histogram plots to check Log(Carat_Weight) distribution

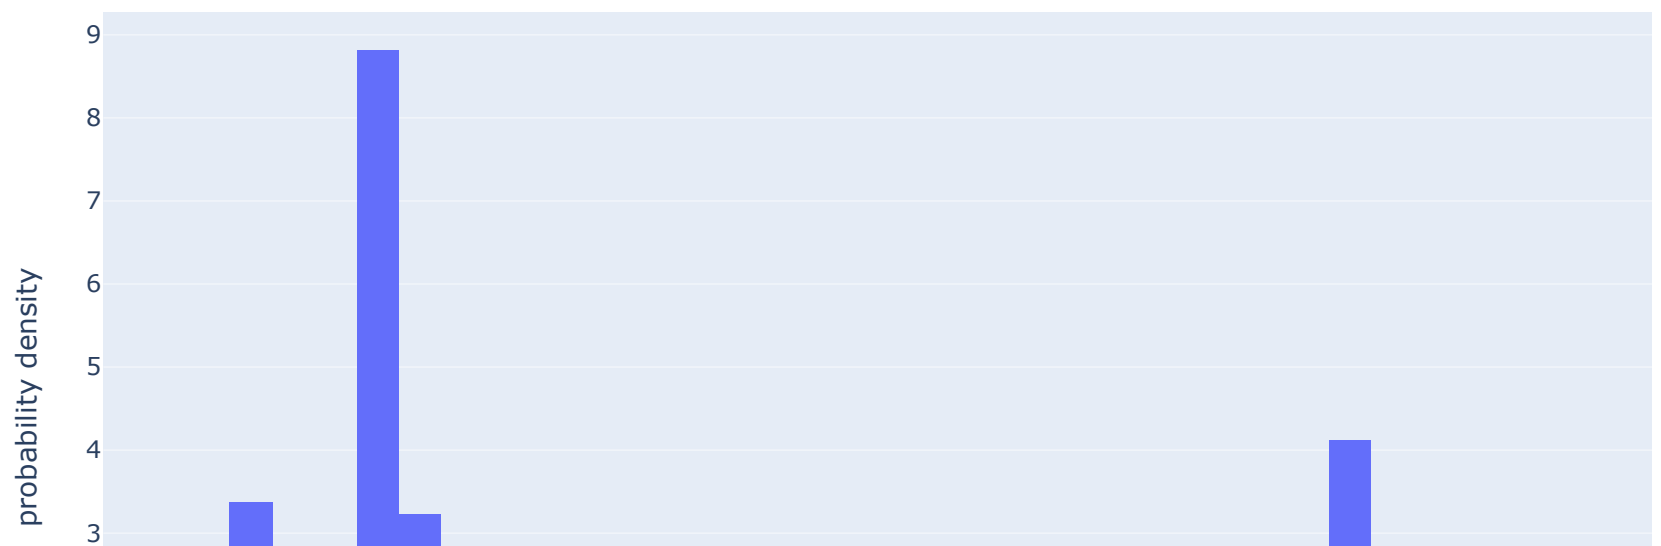
fig = px.histogram(df_reg, x='log_CW', histnorm='probability density')
fig.show()
```





```
In [58]: # histogram plots to check Normalized(Carat_Weight) distribution

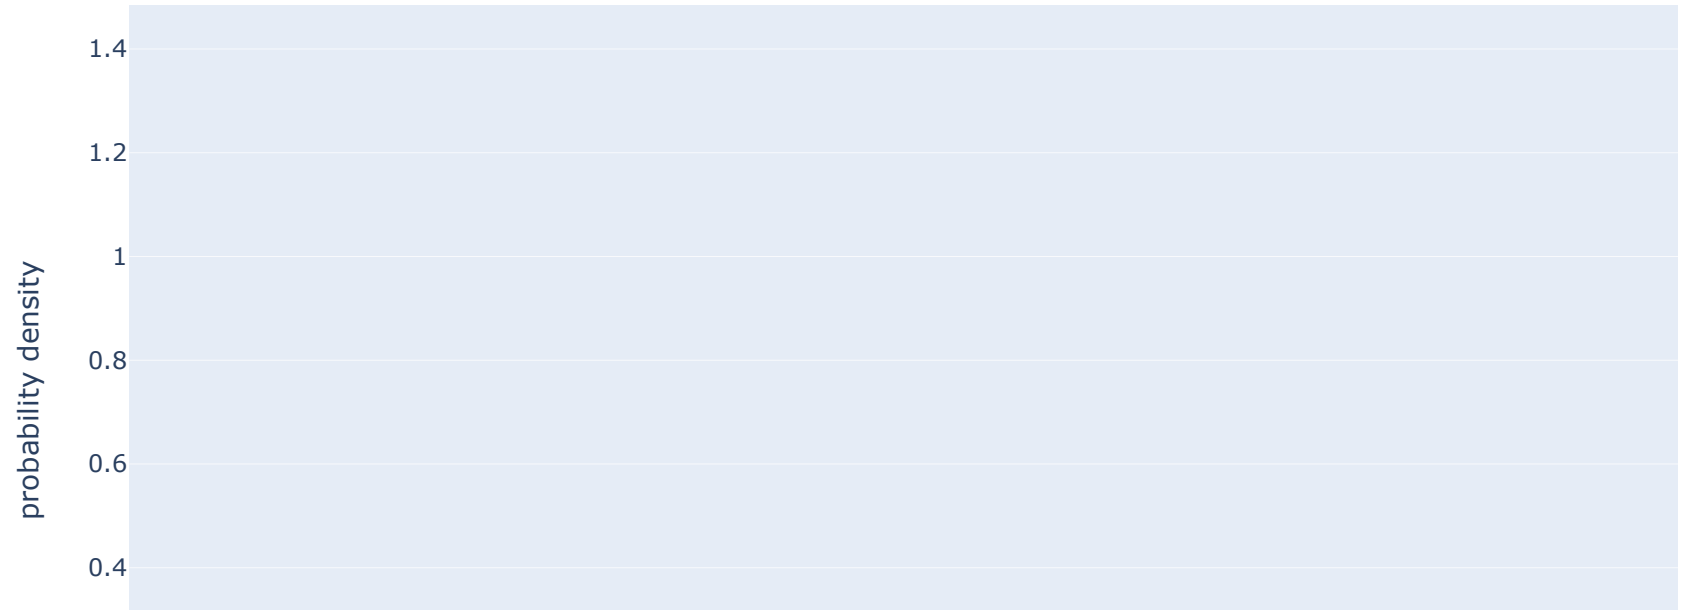
fig = px.histogram(df_reg, x='norm_CW', histnorm='probability density')
fig.show()
```





In [59]:

```
# histogram plots to check Log(Normalized(Carat_Weight)) distribution  
  
fig = px.histogram(df_reg, x='lgnm_CW', histnorm='probability density')  
fig.show()
```



4.2 Pre-processing

```
In [60]: # collecting all the numerical and categorical variables and saving them in two different lists

categorical_columns_seen = [c for i,c in enumerate(df_reg.columns) if df_reg.dtypes[i] in [object]]
numerical_columns_seen = [c for i,c in enumerate(df_reg.columns) if df_reg.dtypes[i] not in [object]]
categorical_columns_seen, numerical_columns_seen
```

```
Out[60]: (['Cut', 'Color', 'Clarity', 'Polish', 'Symmetry', 'Report'],
          ['ID', 'Carat_Weight', 'Price', 'log_CW', 'norm_CW', 'lgnm_CW'])
```

```
In [61]: # creating one hot encoded labels for categorical data

encoder = OneHotEncoder(handle_unknown='error', drop='first')
encoded_df = (pd.DataFrame(encoder.fit_transform(df_reg[categorical_columns_seen]).toarray()).astype(int)
encoded_df.columns = encoder.get_feature_names_out(categorical_columns_seen)
encoded_df.head()
```

```
Out[61]:
```

	Cut_Good	Cut_Ideal	Cut_SignatureIdeal	Cut_Very_Good	Color_E	Color_F	Color_G	Color_H	Color_I	Clarity_IF	...	Clarity_VS2	Clarity_VVS
0	0	1	0	0	0	0	0	1	0	0	...	0	
1	0	1	0	0	0	0	0	1	0	0	...	0	
2	0	1	0	0	0	0	0	1	0	0	...	0	
3	0	1	0	0	1	0	0	0	0	0	...	0	
4	0	1	0	0	0	0	1	0	0	0	...	0	

5 rows × 22 columns



```
In [62]: tmp_df = pd.DataFrame()
tmp_df[numerical_columns_seen] = df_reg[numerical_columns_seen]
```

```
cols = list(encoded_df.columns.values)
tmp_df[cols] = encoded_df[cols]
tmp_df.head()
```

```
Out[62]:
```

	ID	Carat_Weight	Price	log_CW	norm_CW	lgnm_CW	Cut_Good	Cut_Ideal	Cut_SignatureIdeal	Cut_Very_Good	...	Clarity_VS2	Clari
0	1	1.10000	5169.00000	0.09531	0.16204	-1.81987	0	1	0	0	...	0	
1	2	0.83000	3470.00000	-0.18633	0.03704	-3.29557	0	1	0	0	...	0	
2	3	0.85000	3183.00000	-0.16252	0.04630	-3.07248	0	1	0	0	...	0	
3	4	0.91000	4370.00000	-0.09431	0.07407	-2.60255	0	1	0	0	...	0	
4	5	0.83000	3171.00000	-0.18633	0.03704	-3.29557	0	1	0	0	...	0	

5 rows × 28 columns



4.3 Model Development

```
In [63]: # creating Train and test data
Y_reg = tmp_df['Price']
X_reg = tmp_df.drop(['ID', 'Price', 'norm_CW', 'lgnm_CW', 'Carat_Weight'], axis = 1)

X_train_reg, X_test_reg, Y_train_reg, Y_test_reg = train_test_split(X_reg, Y_reg, train_size = 0.8334, random_state = 25)
X_train_reg.shape, X_test_reg.shape
```

```
Out[63]: ((5000, 23), (1000, 23))
```

```
In [64]: # model 1 : Light Gradient Boosting model

lgb_model = lgb.LGBMRegressor(boosting_type='goss', num_leaves=16, max_depth=- 1, learning_rate=0.1, n_estimators=64,
                               objective='regression', reg_lambda=0.4, random_state=25, n_jobs=- 1, importance_type='gain')

lgb_model.fit(X_train_reg, Y_train_reg)
Y_pred_reg = lgb_model.predict(X_test_reg)
```

```
In [65]: # checking model performance on Train and test data
```



```
train_metrics(lgb_model, X_train_reg, Y_train_reg)
test_metrics(Y_test_reg, Y_pred_reg)
```

Cross Validated Metric Results for Train Data:

	R-sq	MAE	MAPE
0	0.98	717.63	6.61
1	0.96	824.34	6.46
2	0.98	760.91	6.45
3	0.98	783.78	6.30
4	0.96	1001.80	7.31
5	0.96	820.40	6.44
6	0.96	833.13	5.94
7	0.96	889.92	6.85
8	0.98	792.09	6.68
9	0.97	825.30	6.84
Mean	0.97	824.93	6.59
Std Dev	0.01	73.47	0.35

Metric Results for Test Data:

	R-sq	MAE	MAPE
0	0.95	810.86	6.38

4.4 Hyperparameter Tuning

In [66]:

```
params = {'n_estimators' : [64, 128, 256], 'max_depth' : [3,5,7], 'learning_rate' : [0.5, 0.1]}
randomized_object = RandomizedSearchCV(estimator=xgb.XGBRegressor(), scoring=make_scorer(neg_mean_absolute_percentage_err),
                                       random_state=25, cv=10, n_jobs=-1, param_distributions=params, verbose=0)
randomized_object.fit(X_train_reg, Y_train_reg)
```

Out[66]:

[illegible]

```

        predictor=None, random_state=None,
        reg_alpha=None, reg_lambda=None,
        scale_pos_weight=None, subsample=None,
        tree_method=None,
        validate_parameters=None,
        verbosity=None),

    n_jobs=-1,
    param_distributions={'learning_rate': [0.5, 0.1],
                        'max_depth': [3, 5, 7],
                        'n_estimators': [64, 128, 256]},
    random_state=25,
    scoring=make_scorer(neg_mean_absolute_percentage_error))

```

In [67]:

```
# tuned hyperparameters
```

```

print('Best Parameters : {}'.format(randomized_object.best_params_))
print('Best MAPE_score : {}'.format(round(randomized_object.best_score_, 4)))
print('Best model : {}'.format(randomized_object.best_estimator_))
CV_Res = pd.concat([pd.DataFrame(randomized_object.cv_results_['params']),
                    pd.DataFrame(randomized_object.cv_results_['mean_test_score'], columns=['MAPE_score'])], axis=1)
CV_Res = CV_Res.sort_values(by='MAPE_score', ascending=False)
print(CV_Res)

```

```
Best Parameters : {'n_estimators': 128, 'max_depth': 7, 'learning_rate': 0.1}
```

```
Best MAPE_score : -5.2092
```

```

Best model : XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                           colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
                           gamma=0, gpu_id=-1, importance_type=None,
                           interaction_constraints='', learning_rate=0.1, max_delta_step=0,
                           max_depth=7, min_child_weight=1, missing=nan,
                           monotone_constraints='()', n_estimators=128, n_jobs=4,
                           num_parallel_tree=1, predictor='auto', random_state=0, reg_alpha=0,
                           reg_lambda=1, scale_pos_weight=1, subsample=1, tree_method='exact',
                           validate_parameters=1, verbosity=None)

```

	n_estimators	max_depth	learning_rate	MAPE_score
3	128	7	0.10000	-5.20919
9	256	7	0.10000	-5.24864
1	256	5	0.10000	-5.28899
7	256	5	0.50000	-5.77052
5	64	5	0.50000	-5.82666
6	128	3	0.50000	-6.33404
0	256	3	0.10000	-6.75647
4	64	3	0.50000	-7.18114

2	128	3	0.10000	-7.53149
8	64	3	0.10000	-9.12262

In [68]:

```
# Tuned model performance on Train and Test data

Y_pred_reg = randomized_object.best_estimator_.predict(X_test_reg)
train_metrics(randomized_object.best_estimator_, X_train_reg, Y_train_reg)
test_metrics(Y_test_reg, Y_pred_reg)
```

Cross Validated Metric Results for Train Data:

	R-sq	MAE	MAPE
0	0.99	613.45	5.14
1	0.98	645.88	5.06
2	0.98	673.20	5.32
3	0.99	693.89	5.19
4	0.97	784.98	5.44
5	0.99	641.57	5.08
6	0.99	649.51	4.90
7	0.97	738.30	5.53
8	0.98	667.26	5.07
9	0.98	686.24	5.14
Mean	0.98	679.43	5.19
Std Dev	0.01	47.81	0.18

Metric Results for Test Data:

	R-sq	MAE	MAPE
0	0.96	674.15	5.00


In [69]:

```
# creating feature importance plot

xgb.plot_importance(randomized_object.best_estimator_, max_num_features=10)
```

Out[69]:

```
<AxesSubplot:title={'center':'Feature importance'}, xlabel='F score', ylabel='Features'>
```




In [70]:

```
# creating features tree plot to explain tuned-model decisions
```

```
tuned_params = randomized_object.best_estimator_.get_booster()
for importance_type in ('weight', 'gain', 'cover', 'total_gain', 'total_cover'):
    print('%s: ' % importance_type, tuned_params.get_score(importance_type=importance_type))

xgb.to_graphviz(randomized_object.best_estimator_, numtrees=0, size='20,20')
```

```
weight: {'log_CW': 3979.0, 'Cut_Good': 251.0, 'Cut_Ideal': 386.0, 'Cut_SignatureIdeal': 186.0, 'Cut_Very_Good': 210.0,
'Color_E': 233.0, 'Color_F': 327.0, 'Color_G': 535.0, 'Color_H': 540.0, 'Color_I': 601.0, 'Clarity_IF': 333.0, 'Clarity_S
I1': 347.0, 'Clarity_VS1': 285.0, 'Clarity_VS2': 326.0, 'Clarity_VVS1': 201.0, 'Clarity_VVS2': 386.0, 'Polish_G': 150.0,
'Polish_ID': 138.0, 'Polish_VG': 343.0, 'Symmetry_G': 200.0, 'Symmetry_ID': 27.0, 'Symmetry_VG': 353.0, 'Report_GIA': 11
1.0}
gain: {'log_CW': 478909536.0, 'Cut_Good': 4592406.5, 'Cut_Ideal': 11844989.0, 'Cut_SignatureIdeal': 19034940.0, 'Cut_Ver
y_Good': 5173736.0, 'Color_E': 14444548.0, 'Color_F': 21241536.0, 'Color_G': 62449032.0, 'Color_H': 112486464.0, 'Color_
I': 126427296.0, 'Clarity_IF': 430294912.0, 'Clarity_SI1': 229794448.0, 'Clarity_VS1': 61524256.0, 'Clarity_VS2': 7511272
0.0, 'Clarity_VVS1': 116088600.0, 'Clarity_VVS2': 55291200.0, 'Polish_G': 2205670.5, 'Polish_ID': 7328479.0, 'Polish_VG':
3530424.25, 'Symmetry_G': 3260049.0, 'Symmetry_ID': 807730.5, 'Symmetry_VG': 4839952.5, 'Report_GIA': 3286022.25}
cover: {'log_CW': 452.32244873046875, 'Cut_Good': 180.37449645996094, 'Cut_Ideal': 258.0777282714844, 'Cut_SignatureIdea
l': 699.6236572265625, 'Cut_Very_Good': 259.4571533203125, 'Color_E': 390.1673889160156, 'Color_F': 378.9541320800781, 'C
olor_G': 306.6728820800781, 'Color_H': 361.75, 'Color_I': 300.1247863769531, 'Clarity_IF': 586.2973022460938, 'Clarity_SI
1': 412.7521667480469, 'Clarity_VS1': 166.25613403320312, 'Clarity_VS2': 240.76380920410156, 'Clarity_VVS1': 645.55224609
375, 'Clarity_VVS2': 378.63470458984375, 'Polish_G': 372.0, 'Polish_ID': 56.565216064453125, 'Polish_VG': 207.57434082031
25, 'Symmetry_G': 449.7850036621094, 'Symmetry_ID': 55.407405853271484, 'Symmetry_VG': 257.2776184082031, 'Report_GIA': 5
58.7927856445312}
total_gain: {'log_CW': 1905581096960.0, 'Cut_Good': 1152694016.0, 'Cut_Ideal': 4572165632.0, 'Cut_SignatureIdeal': 35404
```

```
In [72]: # creating SHAP-explainer plot for all features but for a single data row

shap.force_plot(explainer.expected_value, shap_values[0,:], X_train_reg.iloc[0, :])
```

